

# Threads & Concurrency



## Practice Exercises

- 4.1 Provide three programming examples in which multithreading provides better performance than a single-threaded solution.

**Answer:**

- A web server that services each request in a separate thread
- A parallelized application such as matrix multiplication where various parts of the matrix can be worked on in parallel
- An interactive GUI program such as a debugger where one thread is used to monitor user input, another thread represents the running application, and a third thread monitors performance

- 4.2 Using Amdahl's Law, calculate the speedup gain of an application that has a 60 percent parallel component for (a) two processing cores and (b) four processing cores.

**Answer:**

- With two processing cores we get a speedup of 1.42 times.
- With four processing cores, we get a speedup of 1.82 times.

- 4.3 Does the multithreaded web server described in Section 4.1 exhibit task or data parallelism?

**Answer:**

Data parallelism. Each thread is performing the same task, but on different data.

- 4.4 What are two differences between user-level threads and kernel-level threads? Under what circumstances is one type better than the other?

**Answer:**

- User-level threads are unknown by the kernel, whereas the kernel is aware of kernel threads.

- b. On systems using either many-to-one or many-to-many model mapping, user threads are scheduled by the thread library, and the kernel schedules kernel threads.
- c. Kernel threads need not be associated with a process, whereas every user thread belongs to a process. Kernel threads are generally more expensive to maintain than user threads, as they must be represented with a kernel data structure.

4.5 Describe the actions taken by a kernel to context-switch between kernel-level threads.

**Answer:**

Context switching between kernel threads typically requires saving the value of the CPU registers from the thread being switched out and restoring the CPU registers of the new thread being scheduled.

4.6 What resources are used when a thread is created? How do they differ from those used when a process is created?

**Answer:**

Because a thread is smaller than a process, thread creation typically uses fewer resources than process creation. Creating a process requires allocating a process control block (PCB), a rather large data structure. The PCB includes a memory map, a list of open files, and environment variables. Allocating and managing the memory map is typically the most time-consuming activity. Creating either a user thread or a kernel thread involves allocating a small data structure to hold a register set, stack, and priority.

4.7 Assume that an operating system maps user-level threads to the kernel using the many-to-many model and that the mapping is done through LWPs. Furthermore, the system allows developers to create real-time threads for use in real-time systems. Is it necessary to bind a real-time thread to an LWP? Explain.

**Answer:**

Yes. Timing is crucial to real-time applications. If a thread is marked as real-time but is not bound to an LWP, the thread may have to wait to be attached to an LWP before running. Consider a situation in which a real-time thread is running (is attached to an LWP) and then proceeds to block (must perform I/O, has been preempted by a higher-priority real-time thread, is waiting for a mutual exclusion lock, etc.). While the real-time thread is blocked, the LWP it was attached to is assigned to another thread. When the real-time thread has been scheduled to run again, it must first wait to be attached to an LWP. By binding an LWP to a real-time thread, you are ensuring that the thread will be able to run with minimal delay once it is scheduled.

## Exercises

- 4.8 Under what circumstances does a multithreaded solution using multiple kernel threads provide better performance than a single-threaded solution on a single-processor system?

**Answer:**

When a kernel thread suffers a page fault, another kernel thread can be switched in to use the interleaving time in a useful manner. A single-threaded process, however, will not be capable of performing useful work when a page fault takes place. Therefore, in scenarios where a program might suffer from frequent page faults or might have to wait for other system events, a multithreaded solution would perform better even on a single-processor system.

- 4.9 Which of the following components of program state are shared across threads in a multithreaded process?
- Register values
  - Heap memory
  - Global variables
  - Stack memory

**Answer:**

The threads of a multithreaded process share heap memory and global variables. Each thread has a separate set of register values and a separate stack.

- 4.10 Using Amdahl's Law, calculate the speedup gain for the following applications:
- 40 percent parallel with (a) eight processing cores and (b) sixteen processing cores
  - 67 percent parallel with (a) two processing cores and (b) four processing cores
  - 90 percent parallel with (a) four processing cores and (b) eight processing cores

**Answer:**

- (a) 1.53 (b) 1.60
- (a) 1.19 (b) 1.32
- (a) 3.07 (b) 4.70

- 4.11 Determine if the following problems exhibit task or data parallelism:
- Using a separate thread to generate a thumbnail for each photo in a collection
  - Transposing a matrix in parallel

- A networked application where one thread reads from the network and another writes to the network
- The fork-join array summation application described in Section 4.5.2
- The Grand Central Dispatch system

**Answer:**

- Data parallelism
- Data parallelism
- Task parallelism
- Data parallelism
- Task parallelism

4.12 Consider the following code segment:

```
pid_t pid;

pid = fork();
if (pid == 0) { /* child process */
    fork();
    thread_create( . . . );
}
fork();
```

- How many unique processes are created?
- How many unique threads are created?

**Answer:**

There are six processes and two threads.

4.13 The program shown in Figure 4.23 uses the Pthreads API. What would be the output from the program at `LINE C` and `LINE P`?

**Answer:**

Output at `LINE C` is 5. Output at `LINE P` is 0.

4.14 Consider a multicore system and a multithreaded program written using the many-to-many threading model. Let the number of user-level threads in the program be greater than the number of processing cores in the system. Discuss the performance implications of the following scenarios.

- The number of kernel threads allocated to the program is less than the number of processing cores.
- The number of kernel threads allocated to the program is equal to the number of processing cores.
- The number of kernel threads allocated to the program is greater than the number of processing cores but less than the number of user-level threads.

---

```

#include <pthread.h>
#include <stdio.h>

int value = 0;
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pid_t pid;
    pthread_t tid;
    pthread_attr_t attr;

    pid = fork();

    if (pid == 0) { /* child process */
        pthread_attr_init(&attr);
        pthread_create(&tid,&attr,runner,NULL);
        pthread_join(tid,NULL);
        printf("CHILD: value = %d",value); /* LINE C */
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE P */
    }
}

void *runner(void *param) {
    value = 5;
    pthread_exit(0);
}

```

---

**Figure 4.23** C program for Exercise 4.13.

**Answer:**

When the number of kernel threads is less than the number of processors, then some of the processors will remain idle, since the scheduler maps only kernel threads to processors and not user-level threads to processors. When the number of kernel threads is exactly equal to the number of processors, then it is possible that all of the processors will be utilized simultaneously. However, when a kernel thread blocks inside the kernel (due to a page fault or while invoking system calls), the corresponding processor will remain idle. When there are more kernel threads than processors, a blocked kernel thread can be swapped out in favor of another kernel thread that is ready to execute, thereby increasing the utilization of the multiprocessor system.

