

# Chapter 4:

# Threads & Concurrency



# Chapter 4: Threads & Concurrency

---

Overview

Multicore Programming

Multithreading Models

Thread Libraries

Implicit Threading

Threading Issues

Operating System Examples

# Objectives

---

Identify the basic components of a thread, and contrast threads and processes

Describe the major **benefits** and **significant challenges** of designing **multithreaded processes**

Illustrate different approaches to **implicit threading**, including **thread pools**, **fork-join**, **OpenMP** and **Grand Central Dispatch**

Describe how the Windows and Linux operating systems represent threads

Design multithreaded applications using the Pthreads, Java, and Windows threading **APIs**

# 4.1 Overview

---

A **thread** is a basic unit of CPU utilization; it comprises a thread ID, a program counter, a register set, and a stack.

Most modern applications are **multithreaded**

Threads run within application

Multiple tasks with the application can be implemented by separate threads

- Update display

- Fetch data

- Spell checking

- Answer a network request


Process creation is heavy-weight while **thread creation is light-weight**

Can simplify code, **increase efficiency**

**Kernels are generally multithreaded**

# Multithreaded Server Architecture

---



Remote procedure call

# Benefits

---

**Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces

**Resource Sharing** – threads share resources of process, easier than shared memory or message passing

**Economy** – cheaper than process creation, thread switching lower overhead than context switching

**Scalability** – process can take advantage of multiprocessor architectures

## 4.2 Multicore Programming

---

**Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:

**Dividing activities**

**Balance**

**Data splitting**

**Data dependency**

**Testing and debugging**

**Parallelism** implies a system can perform more than one task simultaneously

**Concurrency** supports more than one task making progress


Single processor / core, scheduler providing concurrency

# Multicore Programming (Cont.)

## Types of parallelism

**Data parallelism** – distributes **subsets** of the same data across multiple cores, **same operation** on each

**Task parallelism** – distributing threads across cores, each thread performing **unique (different) operation**



**Figure 4.5** Data and task parallelism.



# Multicore Programming (Cont.)

---

As # of threads grows, so does architectural support for threading

CPUs have cores as well as *hardware threads*


Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

# Concurrency vs. Parallelism

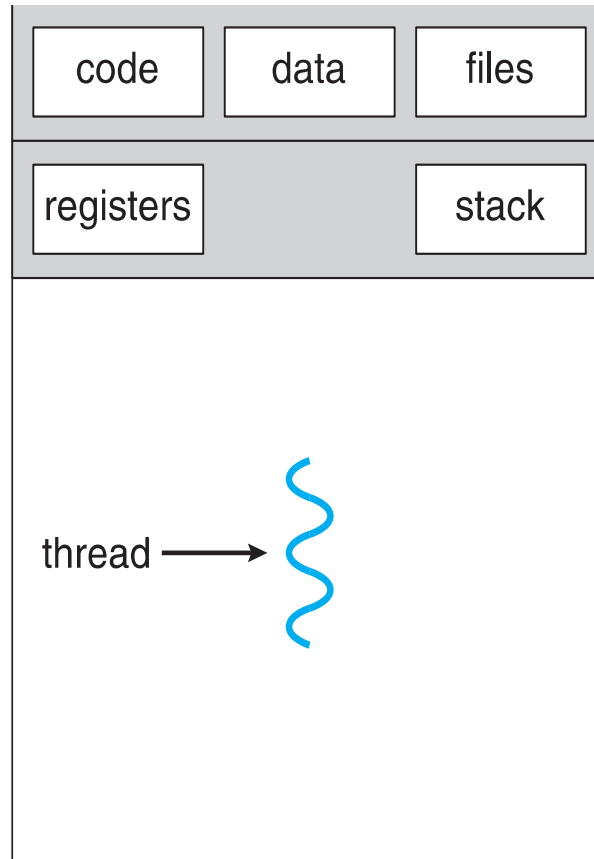
## Concurrent execution on single-core system:



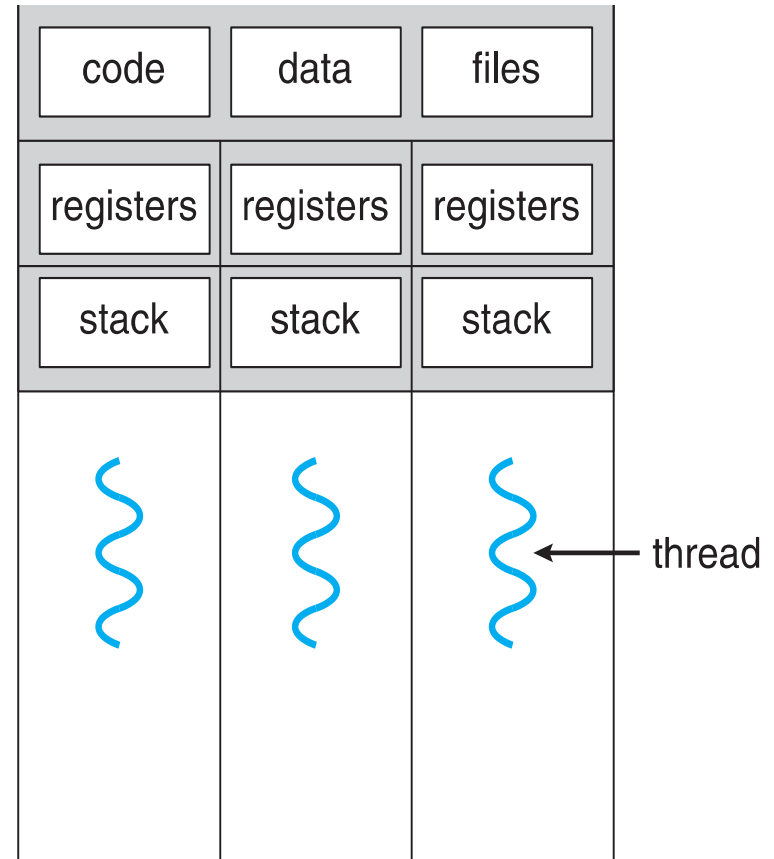
## Parallelism on a multi-core system:



# Single and Multithreaded Processes



single-threaded process



multithreaded process

# Amdahl's Law

---

Identifies performance gains from adding additional cores to an application that has both **serial** and **parallel** components

**S** is serial portion

**N** processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times

As *N* approaches infinity, speedup approaches 1 / *S*

**Serial portion** of an application has **disproportionate effect** on performance gained by adding additional cores

But does the law take into account contemporary multicore systems?

# User Threads and Kernel Threads

**User threads** - management done by **user-level threads library**

Three primary thread libraries:

**POSIX Pthreads**

Windows threads

Java threads

**Kernel threads** - Supported by the Kernel

Examples – virtually all general purpose operating systems, including:


Windows

Solaris

Linux

Tru64 UNIX

Mac OS X



**Figure 4.6** User and kernel threads.

## 4.3 Multithreading Models

---

Many-to-One

One-to-One

Many-to-Many

# Many-to-One

Many **user-level threads** mapped to single **kernel thread**

One thread blocking causes all to block


Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time

Few systems currently use this model

Examples:

**Solaris Green Threads**

**GNU Portable Threads**



# One-to-One

Each user-level thread maps to kernel thread

Creating a user-level thread creates a kernel thread

More concurrency than many-to-one

Number of threads per process sometimes restricted due to **overhead**


Most operating system use this model

Examples

Windows

Linux

Solaris 9 and later





# Many-to-Many Model


Allows many user level threads to be mapped to many kernel threads

Allows the operating system to create a sufficient number of kernel threads

It is difficult to implement

Solaris prior to version 9

Windows with the *ThreadFiber* package



# Two-level Model

Similar to M:M, except that it allows a user thread to be **bound to kernel thread**


Examples

IRIX

HP-UX

Tru64 UNIX

Solaris 8 and earlier



## 4.4 Thread Libraries

---

**Thread library** provides programmer with **API** for creating and managing threads

Two primary ways of implementing

- Library entirely in user space

- Kernel-level library supported by the OS (system call to OS)

# Pthreads

---

May be provided either as user-level or kernel-level

A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization

***Specification***, not ***implementation***

API specifies behavior of the thread library, implementation is up to development of the library

Common in **UNIX operating systems** (Solaris, Linux, Mac OS X)

# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

# Pthreads Example (Cont.)

```
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid,&attr,runner,argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

**Figure 4.9** Multithreaded C program using the Pthreads API.

# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

**Figure 4.10** Pthread code for joining ten threads.

# Windows Multithreaded C Program

---

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```



# Windows Multithreaded C Program (Cont.)

---

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}
```

# Java Threads

---

Java threads are managed by the **JVM**

Typically implemented using the threads model provided by underlying OS

Java threads may be created by:

**Extending Thread class** (derived from Thread class)

Implementing the **Runnable interface**

```
public interface Runnable
{
    public abstract void run();
}
```

# Java Multithreaded Program

---

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```

# Java Multithreaded Program (Cont.)

---

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>"); }
}
```

## 4.5 Implicit Threading

---

Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads

Creation and management of threads done by **compilers** and **run-time libraries** rather than programmers

Four methods explored

- Thread Pools**

- Fork Join**

- OpenMP**

- Grand Central Dispatch (GCD)**

Other methods include Microsoft **Threading Building Blocks (TBB)**,  
`java.util.concurrent` package

# Thread Pools

---

Create a number of threads in a **pool** where they await work

Advantages:

Usually **slightly faster** to service a request with an existing thread than create a new thread

Allows the number of threads in the application(s) to be **bound to the size of the pool**

Separating the task to be performed from mechanics of creating the task allows us to use different strategies for running task

- ▶ i.e. Tasks could be scheduled to run periodically

**Windows API** supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
    * this function runs as a separate thread.  
    */  
}
```

# Fork Join

A **library** manages the number of threads that are created and is also responsible for assigning tasks to threads

In some way, the fork-join model is a **synchronous version of thread pools** in which a library determines the actual number of threads to create





Figure 4.16 Fork-join parallelism.

# Fork Join



**Figure 4.17** Fork-join in Java.



# Fork Join

---

```
ForkJoinPool pool = new ForkJoinPool();  
// array contains the integers to be summed  
int[] array = new int[SIZE];  
  
SumTask task = new SumTask(0, SIZE - 1, array);  
int sum = pool.invoke(task);
```

# Fork Join in Java

---

```
import java.util.concurrent.*;

public class SumTask extends RecursiveTask<Integer>
{
    static final int THRESHOLD = 1000;

    private int begin;
    private int end;
    private int[] array;

    public SumTask(int begin, int end, int[] array) {
        this.begin = begin;
        this.end = end;
        this.array = array;
    }

    protected Integer compute() {
        if (end - begin < THRESHOLD) {
            int sum = 0;
            for (int i = begin; i <= end; i++)
                sum += array[i];

            return sum;
        }
        else {
            int mid = (begin + end) / 2;

            SumTask leftTask = new SumTask(begin, mid, array);
            SumTask rightTask = new SumTask(mid + 1, end, array);

            leftTask.fork();
            rightTask.fork();

            return rightTask.join() + leftTask.join();
        }
    }
}
```

---

Figure 4.18 Fork-join calculation using the Java API.

# OpenMP

Set of **compiler directives** and an **API** for C, C++, FORTRAN

Provides support for **parallel programming** in shared-memory environments

Identifies **parallel regions** – blocks of code that can run in parallel

**#pragma omp parallel**

Create as many threads as there are cores

```
#pragma omp parallel for  
for(i=0;i<N;i++) {  
    c[i] = a[i] + b[i];  
}
```

**Run for loop in parallel**

```
#include <omp.h>  
#include <stdio.h>  
  
int main(int argc, char *argv[])  
{  
    /* sequential code */  
    #pragma omp parallel  
    {  
        printf("I am a parallel region.");  
    }  
  
    /* sequential code */  
  
    return 0;  
}
```

# Grand Central Dispatch(GCD)

---

Apple technology for Mac OS X and iOS operating systems

Extensions to C, C++ languages, API, and run-time library

Allows identification of parallel sections

Manages most of the details of threading

Block is in “^ { }” - `^ { printf("I am a block"); }`

Blocks placed in **dispatch queue**

Assigned to available thread in thread pool when removed from queue

# Grand Central Dispatch(GCD)

---

Two types of dispatch queues:

**serial** – blocks removed in FIFO order, queue is per process, called **main queue**

- ▶ Programmers can create additional serial queues within program

**concurrent** – removed in FIFO order but several may be removed at a time

- ▶ Three system wide queues with priorities **low, default, high**

```
dispatch_queue_t queue = dispatch_get_global_queue  
    (DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);  
  
dispatch_async(queue, ^{ printf("I am a block."); });
```

# 4.6 Threading Issues

---

Semantics of **fork()** and **exec()** system calls

Signal handling

Synchronous and asynchronous

Thread cancellation of target thread

Asynchronous or deferred

Thread-local storage

Scheduler Activations

# Semantics of `fork()` and `exec()`

---

Does `fork()` duplicate only the calling thread or all threads?

Some UNIXes have two versions of `fork`

`Exec()` usually works as normal – replace the running process including all threads

# Signal Handling

---

**Signals** are used in UNIX systems to notify a process that a particular event has occurred.

A **signal handler** is used to process signals

1. Signal is generated by particular event
2. Signal is delivered to a process
3. Signal is handled by one of two signal handlers:
  1. default
  2. user-defined



# Signal Handling

---

Every signal has **default handler** that kernel runs when handling signal

**User-defined signal handler** can override default

For single-threaded, signal delivered to process

Where should a signal be delivered for multi-threaded?

Deliver the signal to **the thread to which the signal applies** (synchronous)

Deliver the signal to **every thread in the process** (<control ><C> signal)

Deliver the signal to **certain threads in the process**

Assign **a specific thread to receive all signals for the process**

# Thread Cancellation

---

Terminating a thread before it has finished

Thread to be canceled is **target thread**

Two general approaches:

**Asynchronous cancellation** terminates the target thread immediately

**Deferred cancellation** allows the target thread to periodically check if it should be cancelled

Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);
```

# Thread Cancellation (Cont.)

---

Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state (**Pthread support three cancellation modes**)

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

If thread has cancellation disabled, cancellation remains **pending** until thread enables it

Default type is deferred

Cancellation only occurs when thread reaches **cancellation point**

- ▶ i.e. `pthread_testcancel()`
- ▶ Then **cleanup handler** is invoked

On Linux systems, thread cancellation is handled through signals

# Thread-Local Storage

---

**Thread-local storage (TLS)** allows each thread to have its own copy of data

Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

**Different from local variables**

Local variables visible only during **single function invocation**

TLS visible across function invocations

Similar to `static` data

TLS is unique to each thread

# Scheduler Activations

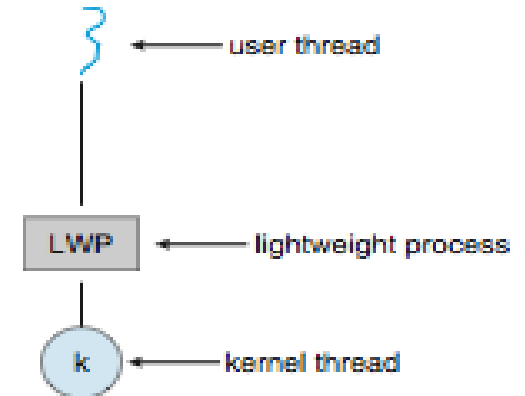
Both M:M and Two-level models require **communication** to maintain the appropriate number of kernel threads allocated to the application

Typically use an **intermediate data structure** between user and kernel threads – **lightweight process (LWP)**

Appears to be a virtual processor on which process can schedule user thread to run

Each LWP attached to kernel thread

How many LWPs to create? (CPU bound : 1 I/O bound : many)




# Scheduler Activations

Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library

This communication allows an application to maintain the correct number kernel threads

5 steps

- ▶ Kernel thread is about to block
- ▶ Upcall with thread ID
- ▶ Kernel **allocates a new virtual processor**
- ▶ Application runs an **upcall handler on it** (save the state of the blocking thread and relinquishes the virtual processor)
- ▶ Upcall handler schedules another thread to run on the new virtual processor



# 4.7 Operating System Examples

---

Windows Threads

Linux Thread

# Windows Threads

---

Windows implements the Windows API – primary API for Win 98, Win NT, Win 2000, Win XP, and Win 7

Implements the **one-to-one mapping**, kernel-level

Each thread contains

- A thread id**

- Register set** representing state of processor

- Separate **user** and **kernel stacks** for when thread runs in user mode or kernel mode

- Private data storage area** used by run-time libraries and dynamic link libraries (DLLs)



# Windows Threads

---

The register set, stacks, and private storage area are known as the **context** of the thread

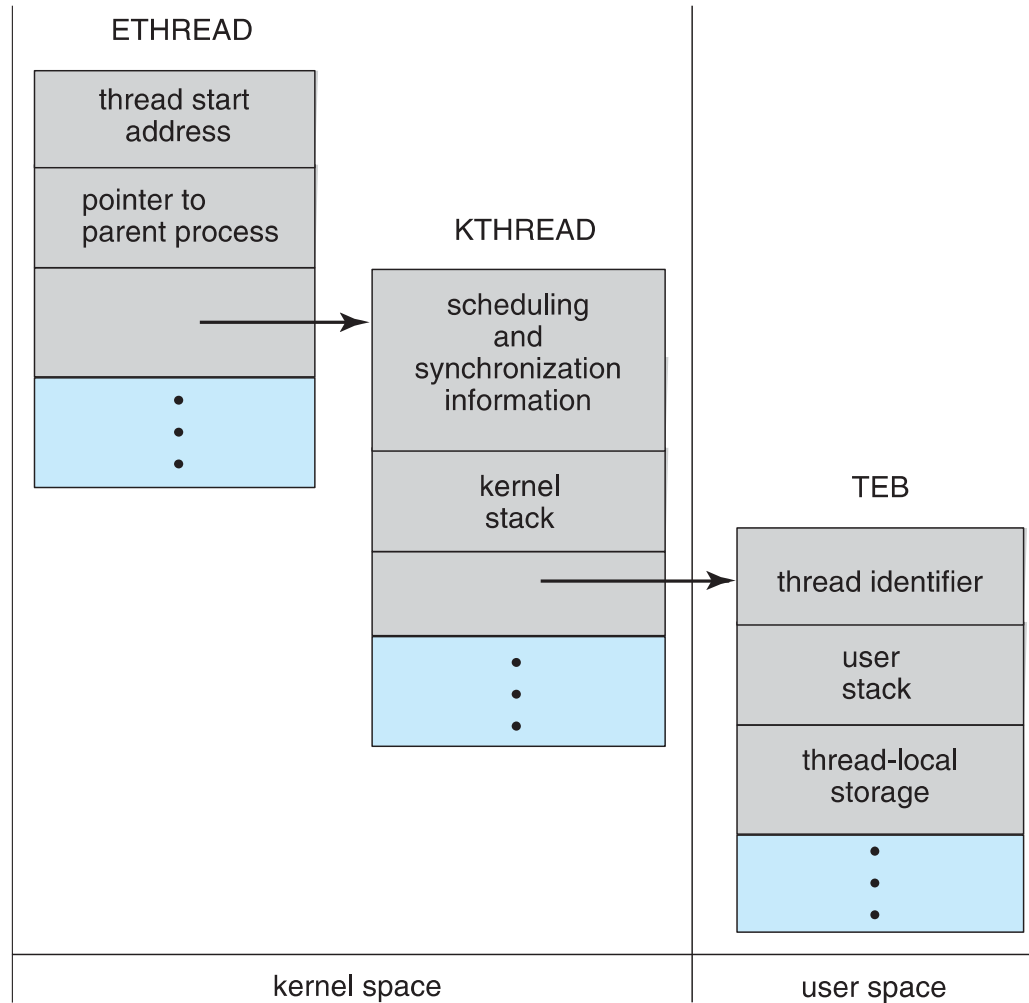
The primary data structures of a thread include:

**TEB (thread environment block)** – thread id, user-mode stack, thread-local storage, in **user space**

**KTHREAD (kernel thread block)** – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in **kernel space**

**ETHREAD (executive thread block)** – includes pointer to process to which thread belongs and to KTHREAD, in **kernel space**

# Windows Threads Data Structures



# Linux Threads

Linux provides the **fork()** system call with the traditional functionality of duplicating a process

Linux refers to a flow of control within a program as **tasks** rather than **threads**

Thread creation is done through **clone()** system call

**clone()** allows a child task to **share** the address space of the parent task (process)

Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

**struct task\_struct** points to process data structures (shared or unique)

Different between **fork()** (copy all) and **clone()** ( may be pointer)

# End of Chapter 4

