

Computer Vision

Ch.8 Image Smoothing

Prof. Po-Yueh Chen (陳伯岳)

E-mail: pychen@cc.ncue.edu.tw

Ext: 8440

NCUE CSIE

Review (1/3)

➤ Noise

Common noise


- Salt-pepper noise
- Gaussian noise
- Spike noise
- Shot noise
- ... etc.




Image with salt and pepper noise

Review (2/3)


➤ Median filter



Window of size 5 in 1D.



Taking the median.




Window of size 3x3 in 2D.

Review (3/3)

➤ Median filter

- ✓ The Median Filter is a non-linear digital filtering technique, often used to remove noise from an image or signal.



Introduction of Image Smoothing

- ✓ Smoothing is also called **blurring**.
 - ✓ It is a simple and frequently used image processing operation.
 - ✓ There are many reasons for smoothing.
 - ✓ It is usually done to **reduce noise or camera artifacts**.
-
- Image smoothing can be done by spatial filter
 - Mean filter
 - Gaussian filter
 - Median filter
 - Bilateral filter

Image Filter – Mean filter (1/3)

➤ Mean filter

- ✓ Mean filter is also called average filter.
- ✓ It smooths an image by replacing each pixel by the average pixel value computed over a rectangular neighborhood (kernel).

When performing smoothing with mean filter:

- Compute the average pixel value over this region.
- Save the result to the anchor pixel.




Image Filter – Mean filter (2/3)

➤ Code

Syntax: ✓ Blurs an image using the normalized box filter.

```
blur(src, dst, ksize, anchor, borderType);
```

src – Input image; it can have any number of channels.

dst – Output image of the same size and type as src.

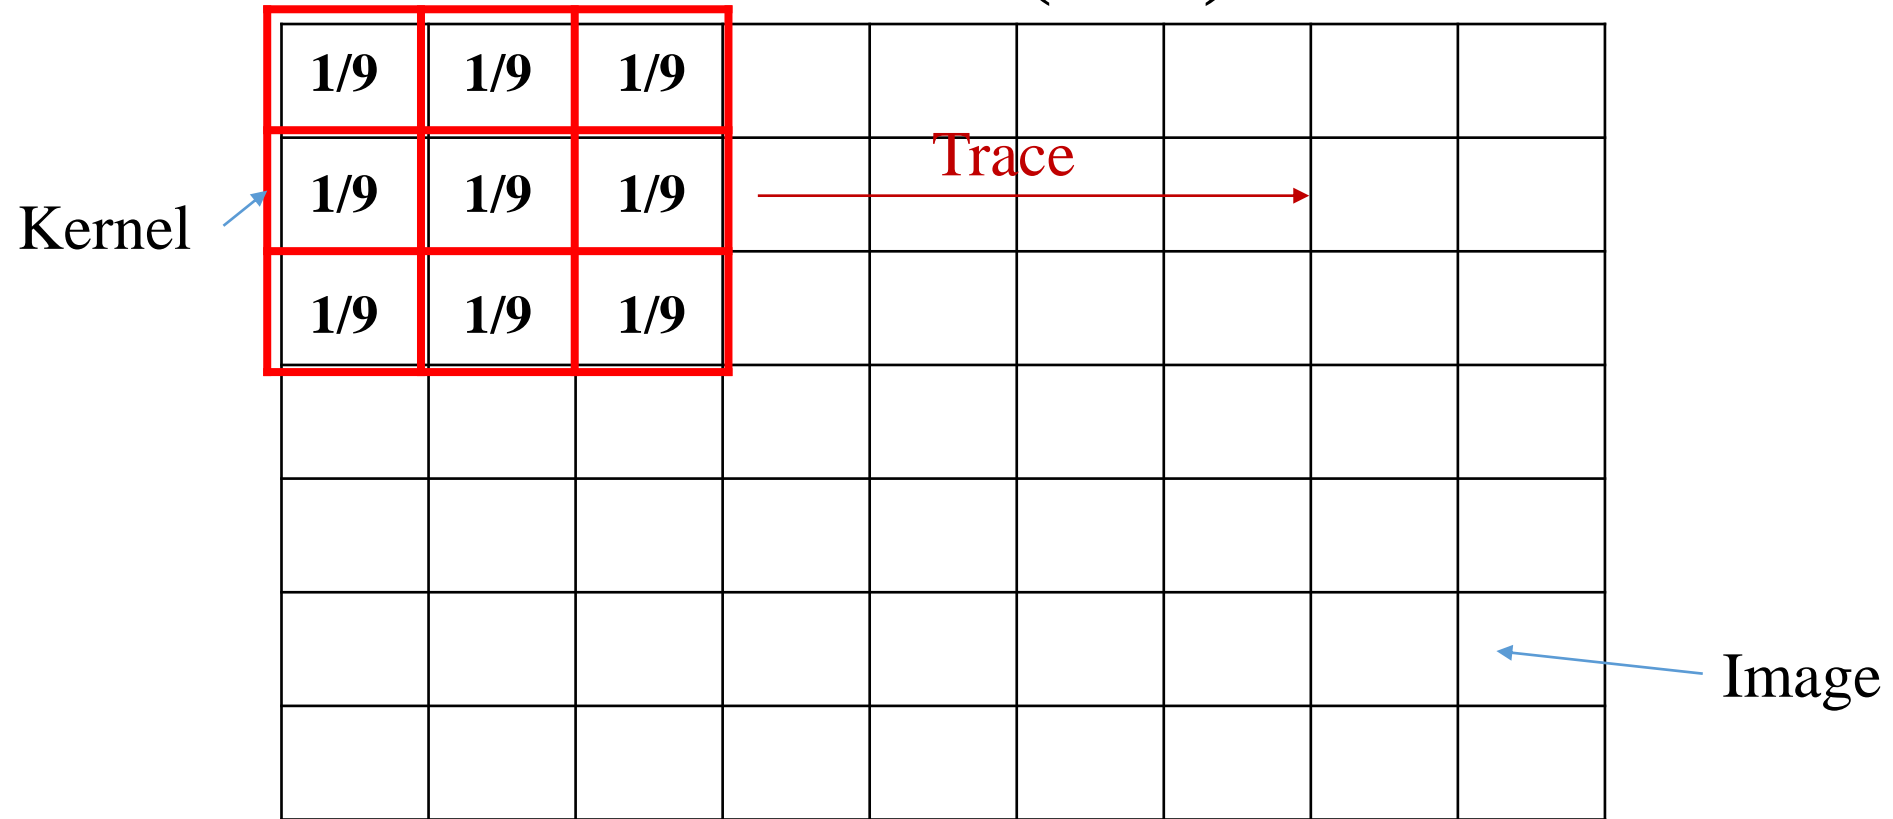
ksize – Blurring kernel size.

anchor – Anchor point; default value Point(-1,-1) means that the anchor is at the kernel center.

borderType – Border mode used to extrapolate pixels outside of the image.

✓ The call `blur(src, dst, ksize, anchor, borderType)` is equivalent to `boxFilter(src, dst, src.type(), anchor, true, borderType)`.

Image Filter – Mean filter (3/3)



- A 3x3 kernel is used to perform mean smoothing.
- It uses such a kernel to scan the whole image.
- At each location, it computes **the average pixel value** under the kernel, and saves the result to the kernel center.

Introduction Gaussian filter & Convolutional

- ✓ The next smoothing filter, Gaussian filter, is probably the most useful.
- ✓ Gaussian filtering **convolves** each point in the image with a Gaussian kernel and then sum to produce the output image.

Two key points to be introduced first:

- What is the operation of convolution?
 - ✓ In AI and Deep Learning, the famous term **CNN** is the abbreviation of **Convolutional Neural Network**.
- What is Gaussian kernel?

Image Filter – Convolution (1/14)

➤ Convolution

- ✓ Convolution is like computing the weighted sum of the pixel values under the kernel.
- ✓ For example, a 3×3 square kernel is used.
- ✓ The kernel coefficients $w(i, j)$ are the weights of pixels $f(x + i, y + j)$ under the kernel.

Image Filter – Convolution (2/14)

➤ Convolution

- ✓ We multiple the pixel value $f(x + i, y + j)$ by the corresponding weight $w(i, j)$, and then sum them up, as the following equation:

$$\begin{aligned} g(x, y) = & w(-1, -1) \times f(x - 1, y - 1) + \\ & w(-1, 0) \times f(x - 1, y) + \\ & w(-1, 1) \times f(x - 1, y + 1) + \\ & w(0, -1) \times f(x, y - 1) + \\ & w(0, 0) \times f(x, y) + \\ & w(0, 1) \times f(x, y + 1) + \\ & w(1, -1) \times f(x + 1, y - 1) + \\ & w(1, 0) \times f(x + 1, y) + \\ & w(1, 1) \times f(x + 1, y + 1) \end{aligned}$$

Image Filter – Convolution (3/14)

➤ Convolution

$$\begin{aligned}
 g(x, y) = & w(-1, -1) \times f(x - 1, y - 1) + \\
 & w(-1, 0) \times f(x - 1, y) + \\
 & w(-1, 1) \times f(x - 1, y + 1) + \\
 & w(0, -1) \times f(x, y - 1) + \\
 & w(0, 0) \times f(x, y) + \\
 & w(0, 1) \times f(x, y + 1) + \\
 & w(1, -1) \times f(x + 1, y - 1) + \\
 & w(1, 0) \times f(x + 1, y) + \\
 & w(1, 1) \times f(x + 1, y + 1)
 \end{aligned}$$


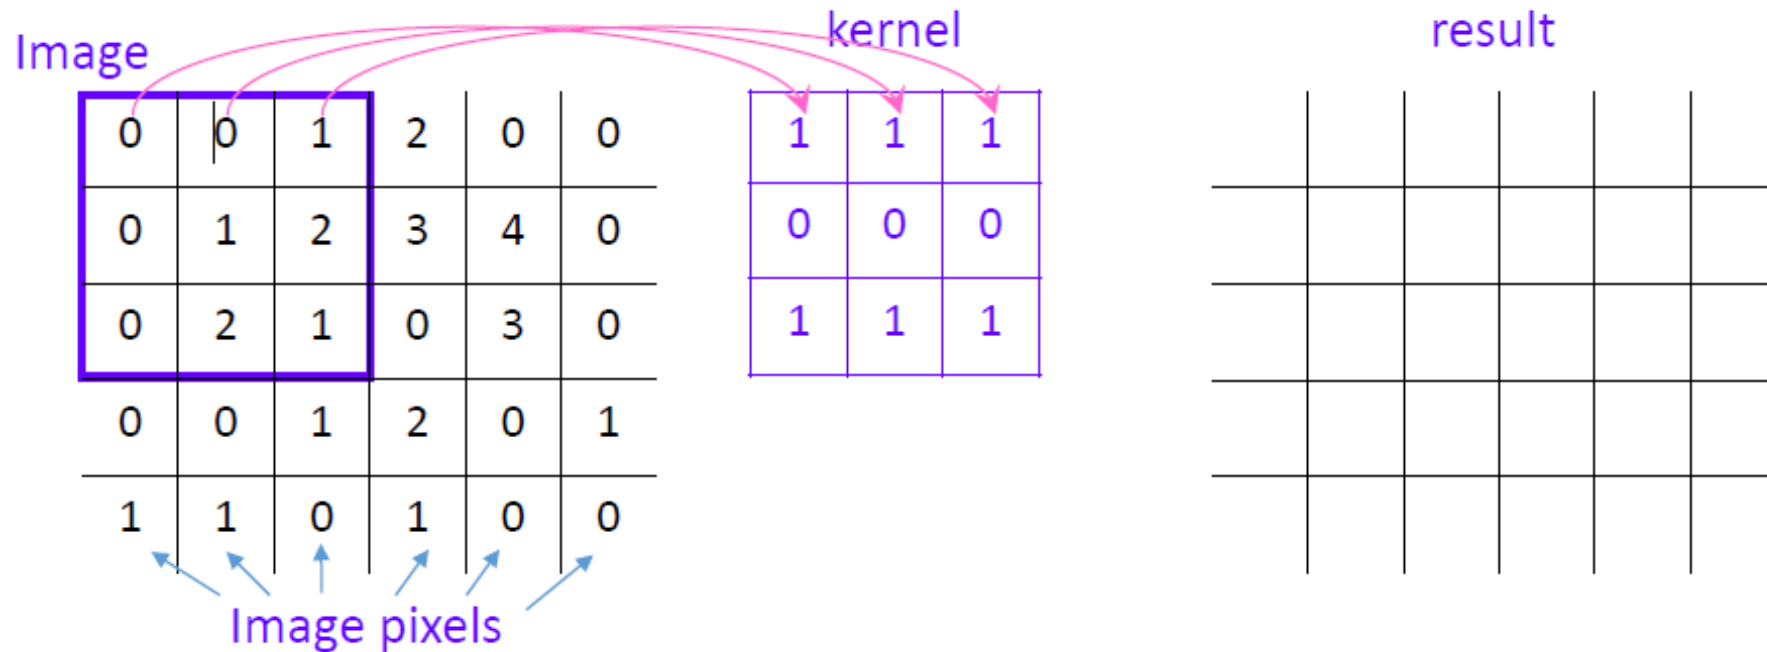


Image Filter – Convolution (4/14)

➤ Convolution

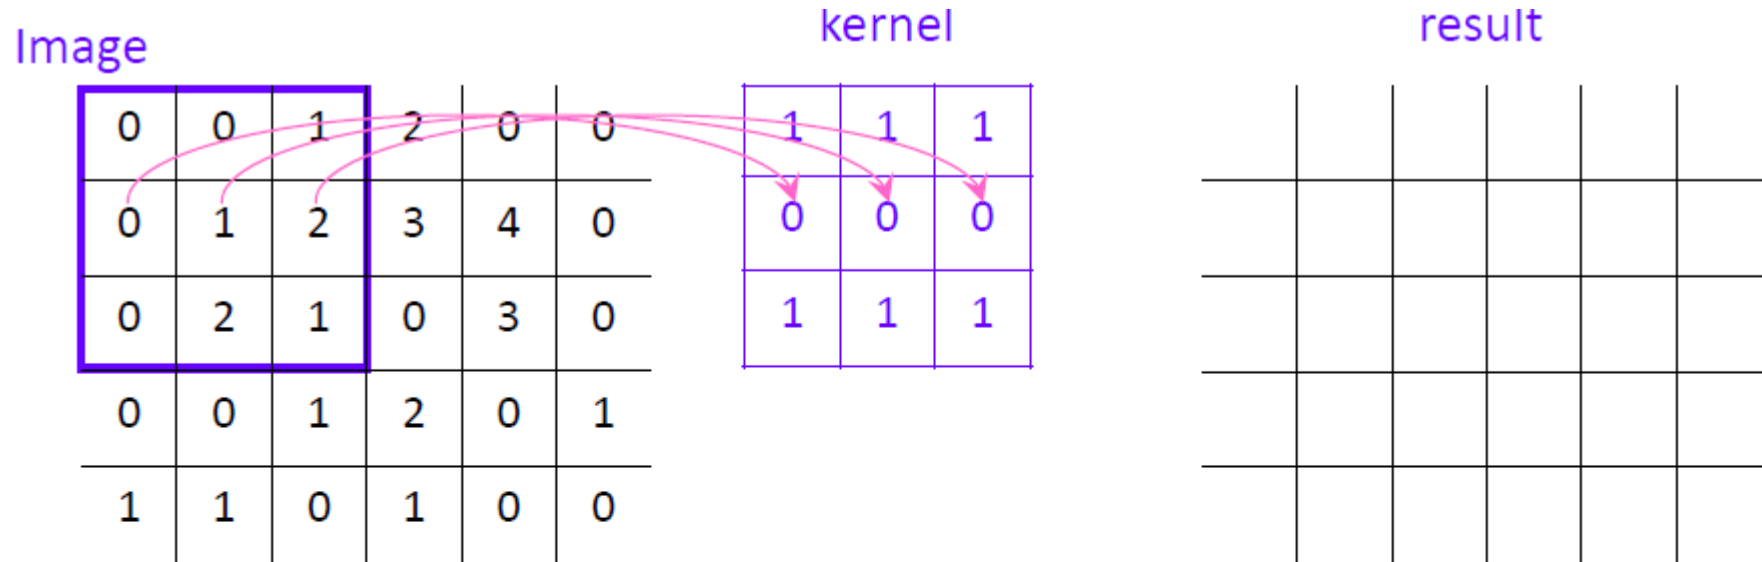


$$\begin{aligned} &\Rightarrow 0 \times 1 + 0 \times 1 + 1 \times 1 + \\ &\quad 0 \times 0 + 1 \times 0 + 2 \times 0 + \\ &\quad 0 \times 1 + 2 \times 1 + 1 \times 1 \\ &= 4 \end{aligned}$$

✓ A simple example is used to explain the basic idea of convolution.

Image Filter – Convolution (5/14)

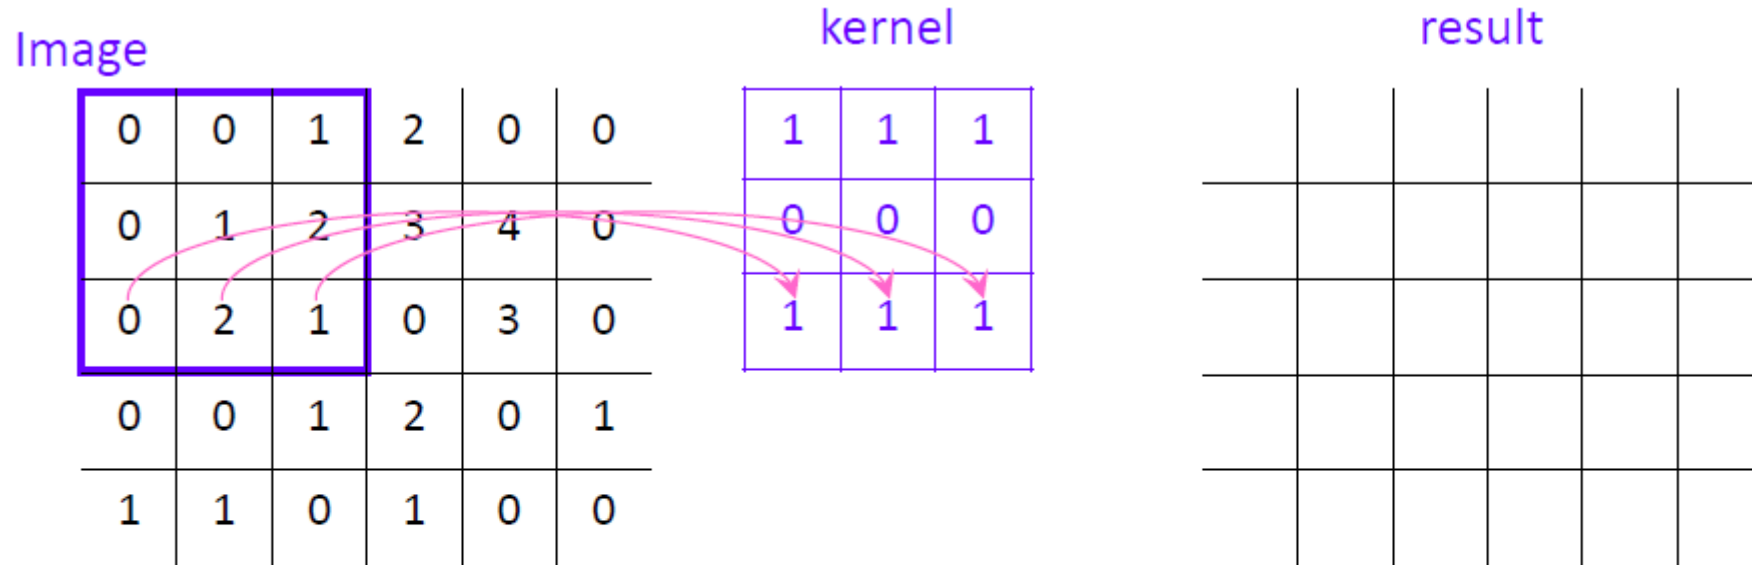
➤ Convolution



$$\begin{aligned} & 0 \times 1 + 0 \times 1 + 1 \times 1 + \\ \Rightarrow & 0 \times 0 + 1 \times 0 + 2 \times 0 + \\ & 0 \times 1 + 2 \times 1 + 1 \times 1 \\ & = 4 \end{aligned}$$

Image Filter – Convolution (6/14)

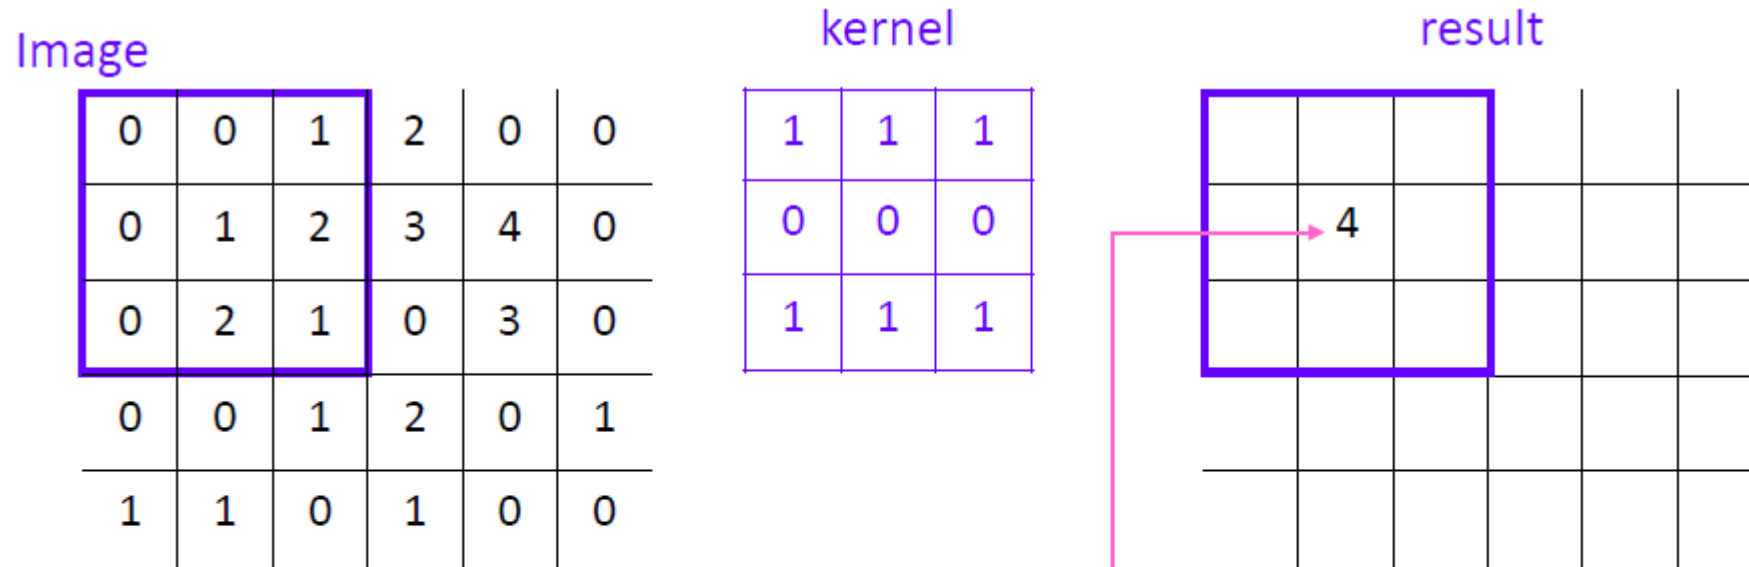
➤ Convolution



$$\begin{aligned} &0 \times 1 + 0 \times 1 + 1 \times 1 + \\ &0 \times 0 + 1 \times 0 + 2 \times 0 + \\ \Rightarrow &0 \times 1 + 2 \times 1 + 1 \times 1 \\ &= 4 \end{aligned}$$

Image Filter – Convolution (7/14)

➤ Convolution



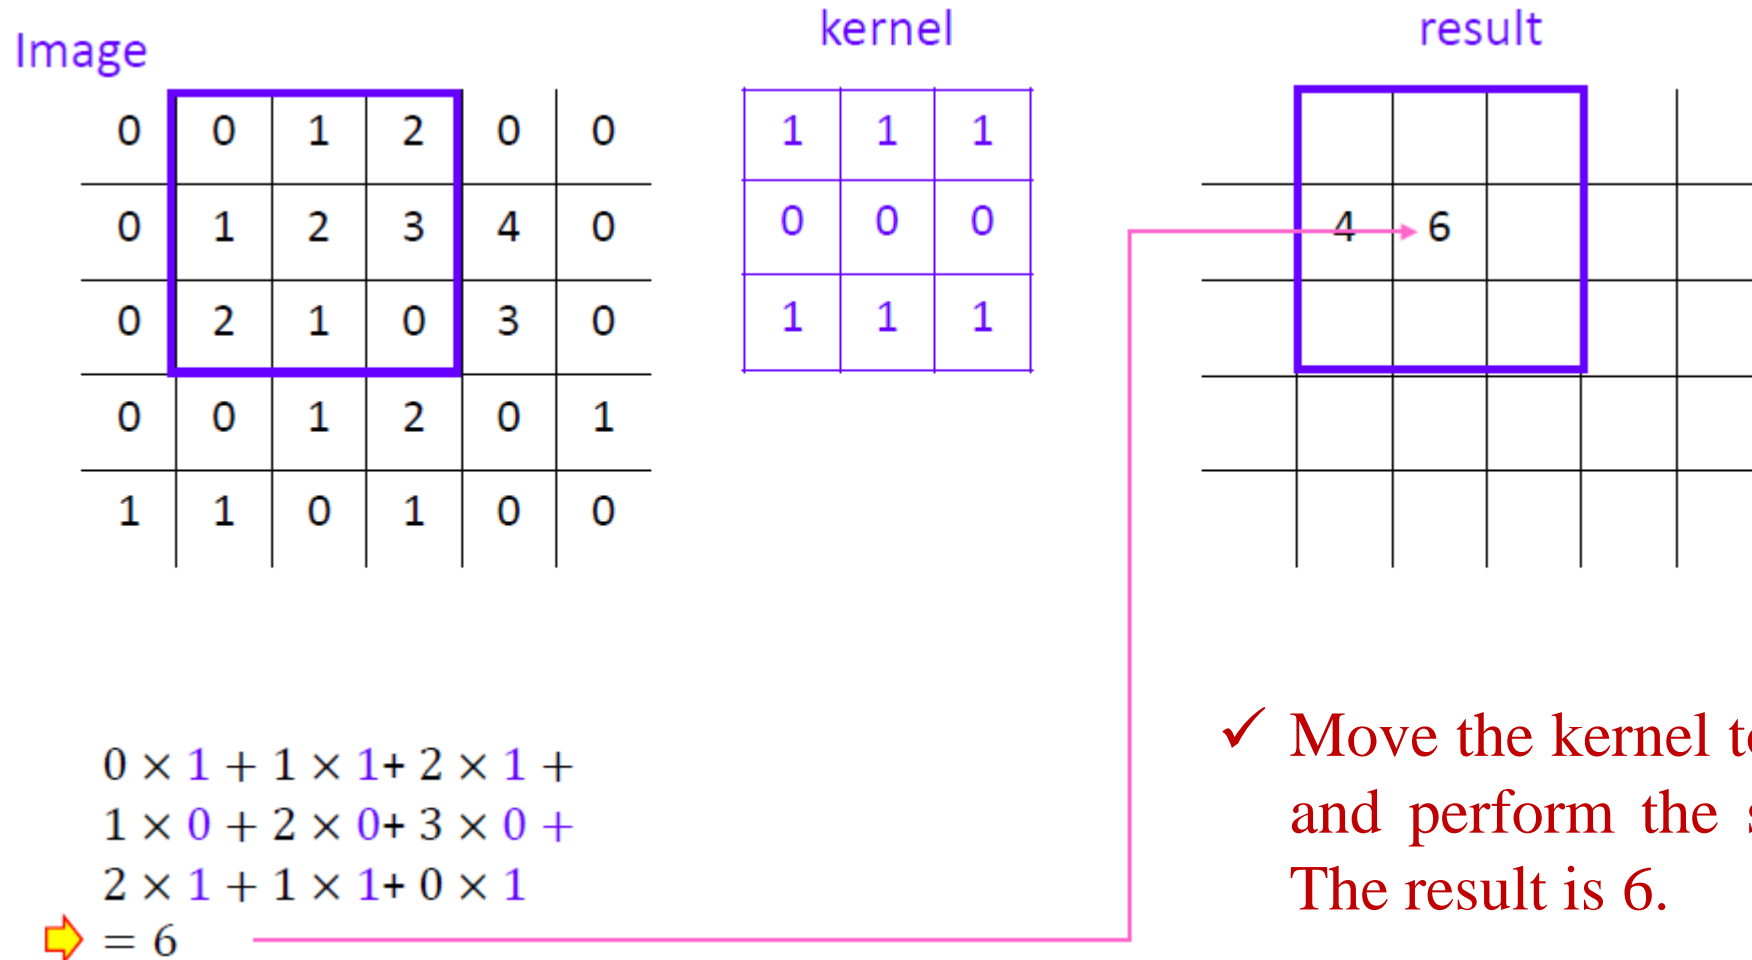
$$\begin{aligned} &0 \times 1 + 0 \times 1 + 1 \times 1 + \\ &0 \times 0 + 1 \times 0 + 2 \times 0 + \\ &0 \times 1 + 2 \times 1 + 1 \times 1 \end{aligned}$$

➡ = 4

✓ To sum them up. Then we obtain the result, 4, and save it to the center pixel.

Image Filter – Convolution (8/14)

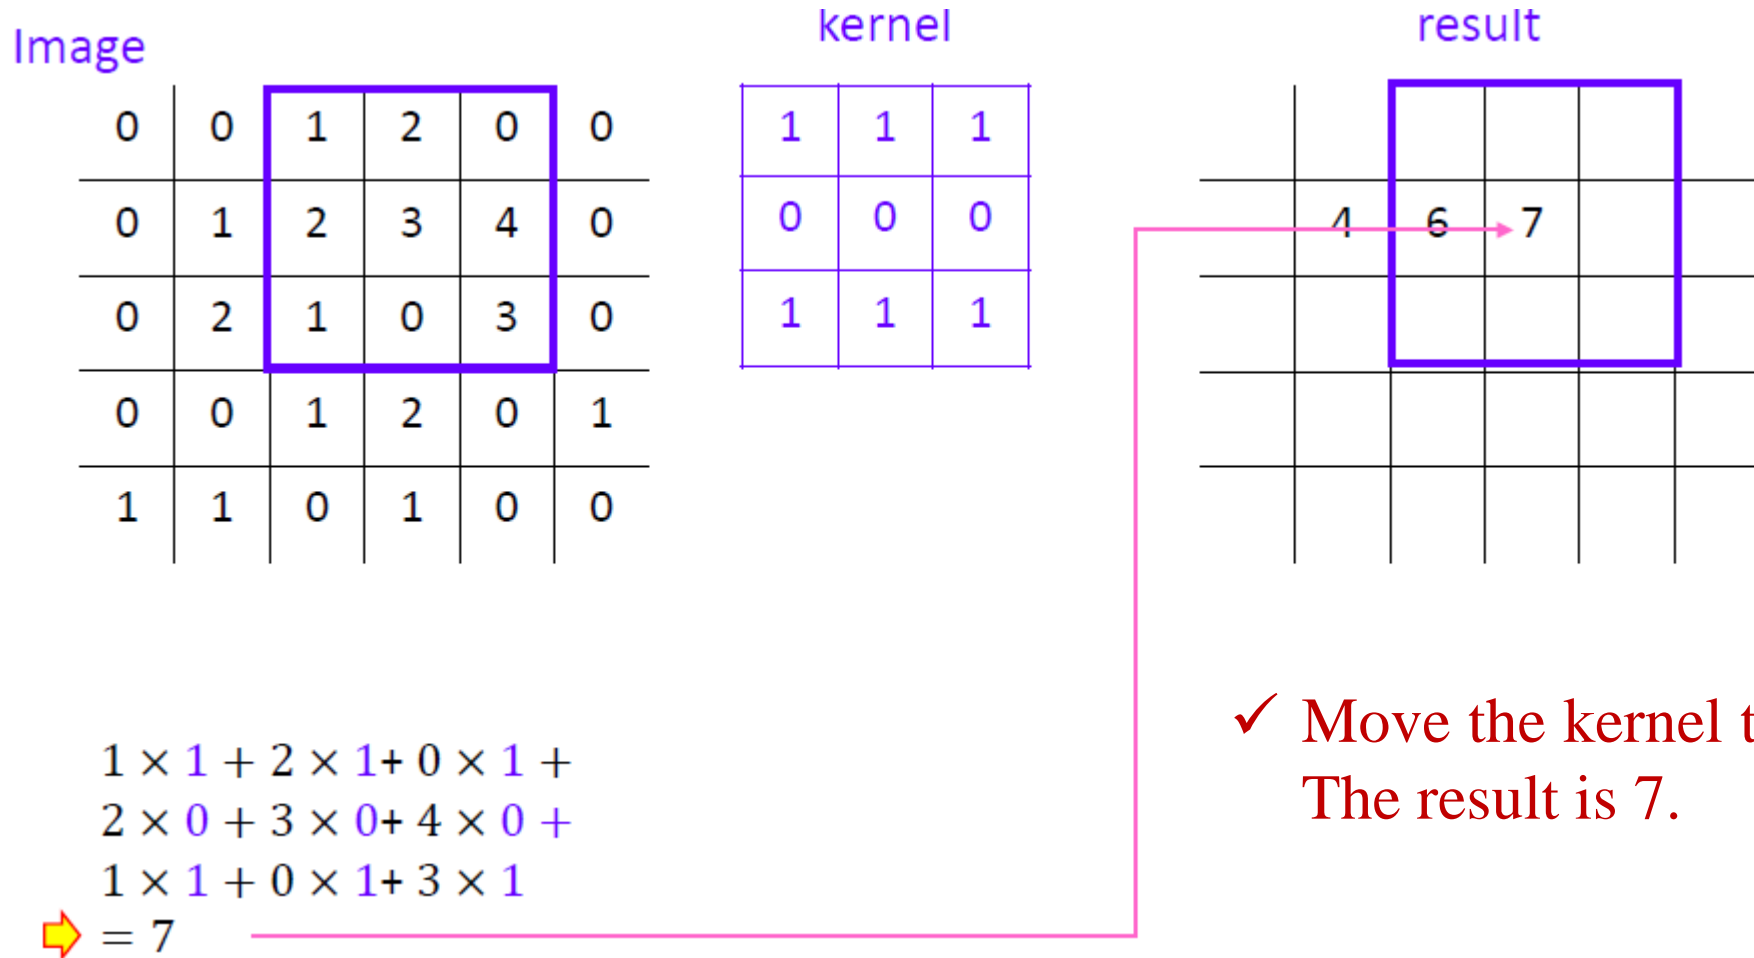
➤ Convolution



✓ Move the kernel to the next position, and perform the same computation. The result is 6.

Image Filter – Convolution (9/14)

➤ Convolution



✓ Move the kernel to the next position.
The result is 7.

Image Filter – Gaussian filter (10/14)

➤ Gaussian kernel

- 2D Gaussian function

$$f(x, y) = A \exp \left(- \left(\frac{(x - x_o)^2}{2\sigma_x^2} + \frac{(y - y_o)^2}{2\sigma_y^2} \right) \right)$$


A is the amplitude, (x_o, y_o) is the center,

σ_x, σ_y are the standard variations in x - and y - directions.

Image Filter – Gaussian filter (11/14)

➤ Gaussian filter

- ✓ A main characteristic of smoothing by Gaussian filter is that **the pixels closer to the center are weighted more.**
- ✓ This is because for the center location, closer pixels are usually more important than farther pixels.

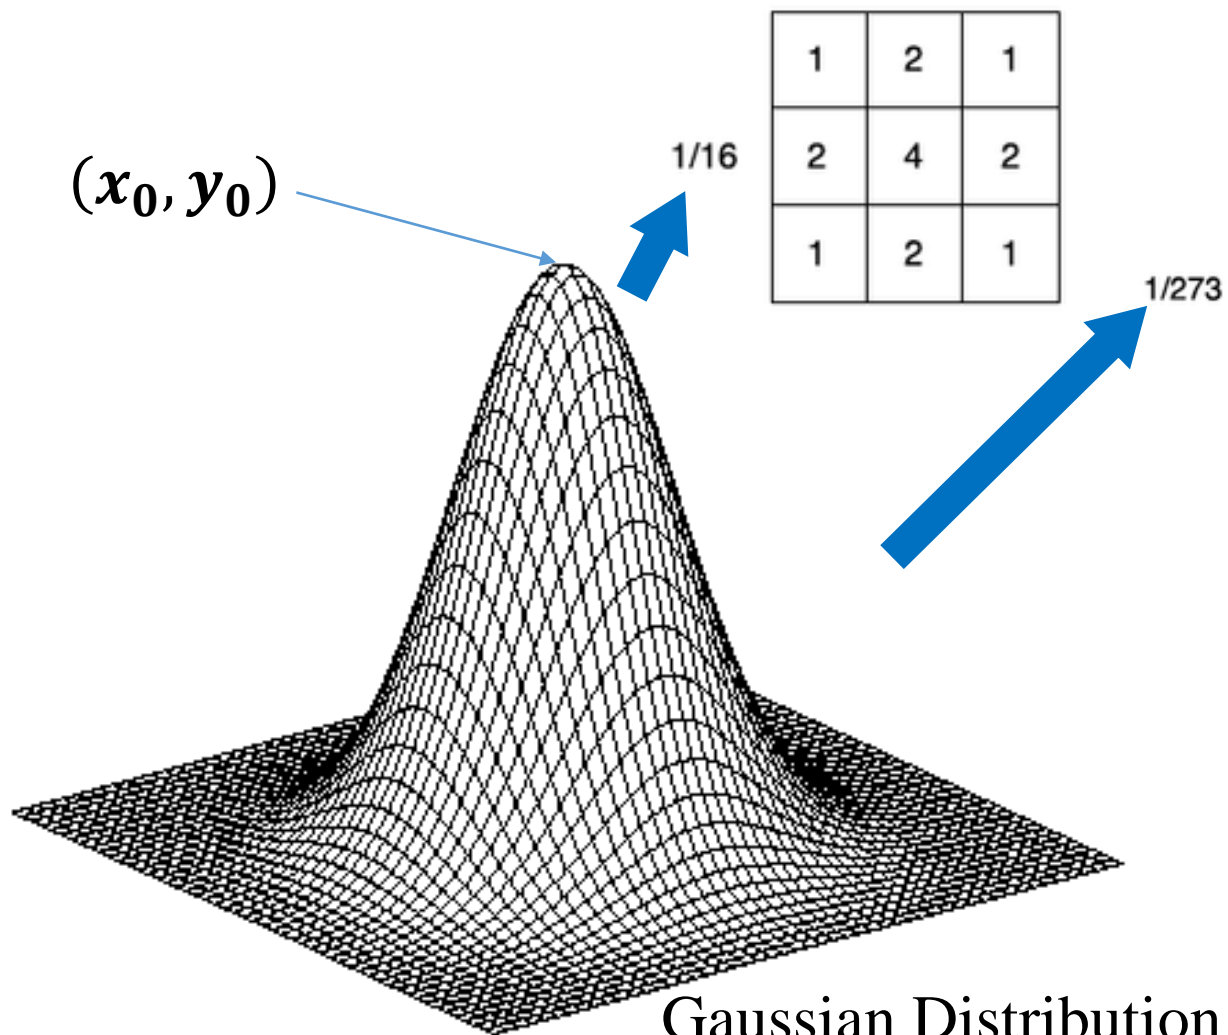


Gaussian Distribution

Image Filter – Gaussian filter (12/14)

➤ Gaussian kernel

✓ From the Gaussian function, we can generate the discrete approximation of Gaussian kernels of different sizes, e.g., 3×3 , 5×5 , 7×7 , and so on.



1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

0	0	1	2	1	0	0
0	3	13	22	13	3	0
1	13	59	97	59	13	1
2	22	97	159	97	22	2
1	13	59	97	59	13	1
0	3	13	22	13	3	0
0	0	1	2	1	0	0

Image Filter – Gaussian filter (13/14)

➤ Gaussian filter

1/1003

0	0	1	2	1	0	0
0	3	13	22	13	3	0
1	13	59	97	59	13	1
2	22	97	159	97	22	2
1	13	59	97	59	13	1
0	3	13	22	13	3	0
0	0	1	2	1	0	0

1/273

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

Gaussian kernel

✓ The pixels closer to the center are weighted more.

Image Filter – Gaussian filter (14/14)

➤ Code

Syntax:

`GaussianBlur(src, dst, ksize, sigmaX, sigmaY, borderType);`

src – Input image; it can have any number of channels.

dst – Output image of the same size and type as src.

ksize – Gaussian kernel size. `ksize.width` and `ksize.height` can differ but they both must be positive and odd. Or, they can be zero's and then they are computed from `sigma*`.

sigmaX – Gaussian kernel standard deviation in X direction.

sigmaY – Gaussian kernel standard deviation in Y direction.


borderType – Pixel extrapolation method (see `borderInterpolate()` for details).

Image Filter – Comparison (1/3)

➤ Gaussian filter

Gaussian smoothing **reduces noise** while preserving signal.

Although the image is blurred, the textures are not lost.




Original Image

Gaussian Results

Image Filter – Comparison (2/3)

➤ Mean filter

Detailed textures are lost.




Original Image

Mean Results

Image Filter – Comparison (3/3)

➤ Median filter

There are several areas of semi-equal color.



Original Image

Median Results


Applications

Here is an example that uses the operations of **blurring** and **thresholding** for region extraction. The objective is to eliminate **irrelevant detail** in the image.

“Irrelevant” refers to pixel regions that are **small** compared to kernel size.



(a) Telescope image



(b) Result of Gaussian filtering



(c) Result after thresholding

Image Filter – Bilateral filter (1/7)

- ✓ Bilateral filtering is known as edge-preserving smoothing.



Original Image



Bilateral Results

Image Filter – Bilateral filter (2/7)

A typical motivation for Gaussian smoothing:

- Pixels in a real image should vary slowly over space and thus be correlated to their neighbors.
- Random noise can be expected to vary greatly from one pixel to the next (i.e., noise is spatially **uncorrelated**).

It is in this sense that Gaussian smoothing **reduces noise** while **preserving signal**.

Unfortunately, this method **breaks down near edges**, where you do expect pixels to be **uncorrelated** with their neighbors.

→ As a result, Gaussian smoothing blurs away edges.

✓ Bilateral filtering provides a mean of smoothing an image without blurring away its edges.


Image Filter – Bilateral filter (3/7)

Like Gaussian smoothing, bilateral filtering constructs a **weighted average** of each pixel and its **neighboring pixels**.

The weighting has two components:

- The first is the same weighting used by Gaussian smoothing, based on the **spatial distance** from the center pixel. → **closer pixels are weighted more**.
 - The second is also a Gaussian weighting but is based on the **difference in intensity** from the center pixel. → **similar pixels are weighted more**.
-
- You can think of bilateral filtering as Gaussian smoothing that weights **similar pixels** more, than less similar ones.
 - The effect of this filter is typically to turn an image into what appears to be a **watercolor painting** of the same scene.
 - This can be useful as an aid to **segmenting an image**.

Image Filter – Bilateral filter (4/7)



✓ Edges are preserved.

Image Filter – Bilateral filter (5/7)

➤ Code

Syntax:

```
bilateralFilter(src, dst, d, sigmaColor, sigmaSpace, borderType);
```

src – Source 8-bit or floating-point, 1-channel or 3-channel image.

dst – Destination image of the same size and type as src.

d – Diameter of each pixel neighborhood that is used during filtering. If it is non-positive, it is computed from sigmaSpace.

sigmaColor – Filter sigma in the color space.

sigmaSpace – Filter sigma in the coordinate space.

borderType – border mode used to extrapolate pixels outside of the image, see BorderTypes

Image Filter – Bilateral filter (6/7)



✓ There are areas of semi-equal color.



Original Image



Bilateral filter

Image Filter – Bilateral filter (7/7)



Original Image



Bilateral filter


✓ Larger sigma results in larger areas of semi-equal color.



Original Image



Bilateral filter



✓ Bilateral filter make the image look like a **watercolor painting**.

Thanks!

Any questions?